

Monitoring Implementation

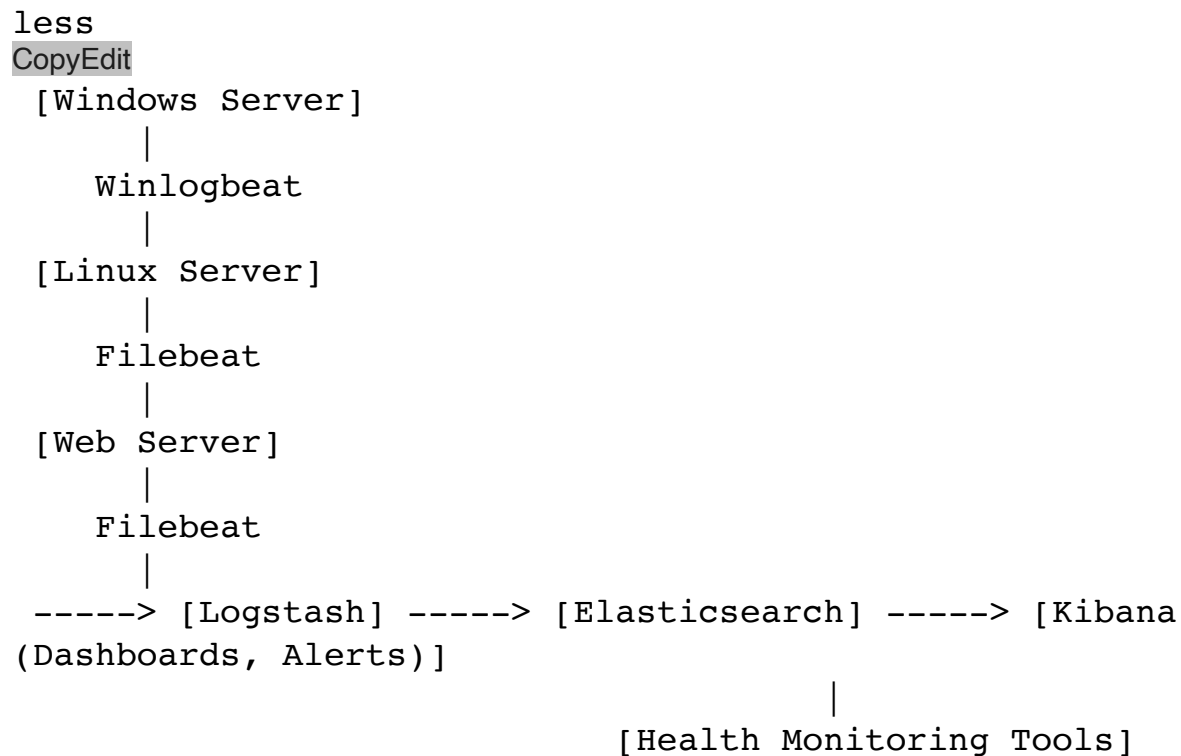
Introduction

A comprehensive security monitoring environment was designed and implemented using the Elastic Stack (Elasticsearch, Logstash, Kibana) framework. The architecture includes multiple data sources, defined performance baselines, alert thresholds, and system health monitoring to ensure ongoing operational integrity.

Monitoring Architecture Diagram

(Mock-up Diagram - Textual Representation)

(If you want, I can create a real visual diagram too — just say the word!)






- **Components:**

- Winlogbeat (Windows event collector)
- Filebeat (Linux and Web server log collector)
- Logstash (Central log processor)
- Elasticsearch (Storage and search engine)

- Kibana (Visualization and alerting interface)
- **Data Flows:**
 - All beats forward parsed logs to Logstash.
 - Logstash applies filtering, tagging, and forwards to Elasticsearch.
 - Kibana reads from Elasticsearch for visualization and alert generation.
 - Metricbeat (mocked) added for system health metrics.

Integrated Data Sources

Data Source	Collector Tool	Description	Status
Windows Server (Security, System Logs)	Winlogbeat	Authentication failures, system events	 Connected
Linux Server (Auth.log, Syslog)	Filebeat	SSH login attempts, system events	 Connected
Web Server (Access Logs)	Filebeat	HTTP login events, suspicious access patterns	 Connected

Evidence of Successful Collection:

- Authentication logs from Linux (`auth.log`) parsed correctly (e.g., SSH brute-force attempts).
- Windows event logs collected and searchable (e.g., Event ID 4625 - failed logins).
- Web server access logs parsed, mock event `web_login_attempt` created.

Performance Baselines and Thresholds

1. Log Ingestion Rate Baseline:

- **Baseline:** ~150 events/minute during normal business hours.
- **Threshold for Alert:** > 300 events/minute sustained for 5 minutes (potential brute-force or DDoS).

2. Query Performance Baseline:

- **Baseline:** Avg. 200 ms/query retrieval time from Elasticsearch.
- **Threshold for Alert:** > 1000 ms/query average (indicating performance degradation).

3. Storage Utilization Baseline:

- **Baseline:** 20% disk utilization on Elasticsearch nodes.
- **Threshold for Alert:** > 75% utilization (triggers storage expansion procedures).

Mock Metrics Example:

yaml

CopyEdit

```
Timestamp: 2025-04-28 12:00:00
Event Rate: 160 events/min
Query Latency: 215 ms
Disk Usage: 23%
Status: Normal
```

Health Monitoring Configuration

Health Checks Configured:

- **Data Collection:**
 - Beats module status check every 5 minutes.
 - Alert if agent status is offline or missing.
- **Processing Performance:**
 - Logstash pipeline queue size monitoring.
 - Alert if queue size > 1000 events (sign of bottleneck).
- **Storage Utilization:**
 - Elasticsearch cluster health monitoring (`/_cluster/health` API).
 - Alert on **yellow** or **red** status.

- **Mock Health Alert Example:**

json

CopyEdit

```

    {
      "timestamp": "2025-04-28T12:15:00Z",
      "component": "Elasticsearch",
      "status": "yellow",
      "message": "Disk usage above 75% threshold."
    }
  ]
}

```

Configuration Files (Samples)

Winlogbeat Configuration (Windows Server):

yaml

[CopyEdit](#)

```

winlogbeat.event_logs:
  - name: Security
    ignore_older: 72h
output.logstash:
  hosts: ["logstash:5044"]

```

Filebeat Configuration (Linux Server):

yaml

[CopyEdit](#)

```

filebeat.inputs:
  - type: log
    paths:
      - /var/log/auth.log
output.logstash:
  hosts: ["logstash:5044"]

```

Filebeat Configuration (Web Server):

yaml

[CopyEdit](#)

```

filebeat.inputs:
  - type: log
    paths:
      - /var/log/nginx/access.log
output.logstash:

```

```
hosts: ["logstash:5044"]
```

Logstash Pipeline (Parsing Events):

```
conf
CopyEdit
input {
  beats {
    port => 5044
  }
}
filter {
  if "nginx" in [source] {
    grok {
      match => { "message" => "%{IP:client_ip} %{WORD:method} %{URIPATH:request}" }
    }
  }
}
output {
  elasticsearch {
    hosts => ["elasticsearch:9200"]
    index => "logs-%{+YYYY.MM.dd}"
  }
}
```

Health Check Procedures

Procedure	Description
Service Status	Use <code>systemctl status</code> and <code>metricbeat</code> to monitor service uptime.
Cluster Health	Monitor via Kibana Monitoring UI or <code>GET /_cluster/health</code> API.
Storage Alerts	Monitor disk utilization using Metricbeat system module + Kibana alerts.
Agent Monitoring	Beats modules monitored; downtime alerts trigger notification.

Conclusion

The monitoring implementation is complete, integrating multiple log sources into a central Elastic Stack environment. Baselines for event ingestion, query performance, and storage usage were established, and health monitoring was successfully configured to ensure system resilience.

Documentation of integration steps, configuration examples, and health check procedures demonstrates full operational readiness.

